



# Overview of the Relational data model

Goals for

Today & Next  
week

(SQL, SQL, SQL)

## Phase I: Intuition for SQL (1st half of today)

Basic Relational model (aka tables)

SQL concepts we'll study (similar to Python map-reduce)

Example SQL (exploring real datasets)

## Phase II: Formal description

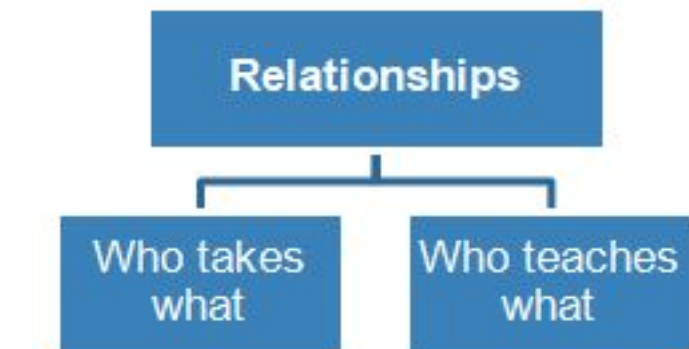
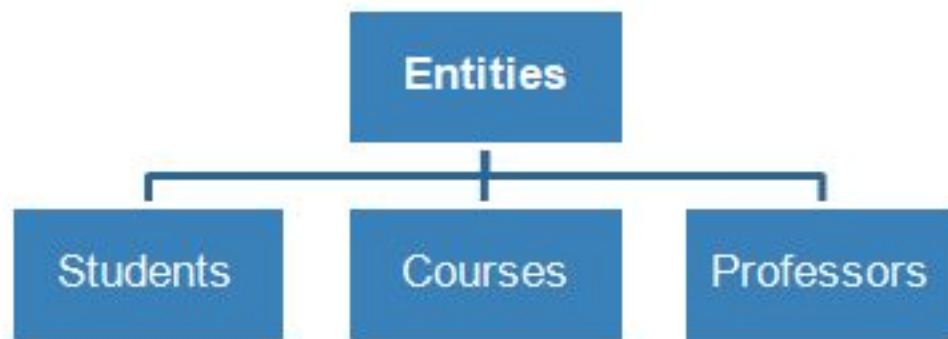
Schemas, Query structure of SELECT-FROM-WHERE, JOINS, etc

# A Motivating Example

A basic Course Management System (CMS):

*Entities* (e.g., Students, Courses)

*Relationships* (e.g., Alice is enrolled in 145)



# Example Spreadsheet Tables

The screenshot shows a Google Sheets spreadsheet with the following tables:

| SID   | Name   | GPA |
|-------|--------|-----|
| 40001 | Mickey | 3.2 |
| 40002 | Daffy  | 3.6 |
| 50003 | Donald | 3.3 |
| 50004 | Minnie | 3.9 |
| 10008 | Pluto  | 4   |

| CID  | C-Name                    | Room       |
|------|---------------------------|------------|
| 1012 | CS145 - Toon systems      | Nvidia     |
| 1017 | CS161 - Animation art     | Gates 300  |
| 1019 | CS245 - Painting town red | Packard 45 |

| SID   | CID  | Grade |
|-------|------|-------|
| 40001 | 1012 | A     |
| 50004 | 1012 | A-    |
| 40001 | 1017 | B+    |

## Logical Schema

Student(sid: *string*, name: *string*, gpa: *float*)

Courses(cid: *string*, c-name: *string*, room: *string*)

Enrolled(sid: *string*, cid: *string*, grade: *string*)

Keys [connect tables]

sid: Connects Enrolled to Student

cid: Connects Enrolled to Courses

## Queries ["compute" over tables]

- Minnie's GPA?
- AVG student GPA?
- Mickey's classes?
- AVG student GPA in CS145?

# Example Spreadsheet Tables

Students

File Edit View Insert Format Data Tools Add-ons Help All changes saved in Drive

100% \$ % .0 .00 123 Arial 10 B I S A

|    | C               | D                         | E           | F | G | H | I |
|----|-----------------|---------------------------|-------------|---|---|---|---|
| 1  |                 |                           |             |   |   |   |   |
| 2  |                 |                           |             |   |   |   |   |
| 3  |                 |                           |             |   |   |   |   |
| 4  | <b>Students</b> |                           |             |   |   |   |   |
| 5  | <b>SID</b>      | <b>Name</b>               | <b>GPA</b>  |   |   |   |   |
| 6  | 40001           | Mickey                    | 3.2         |   |   |   |   |
| 7  | 40002           | Daffy                     | 3.6         |   |   |   |   |
| 8  | 50003           | Donald                    | 3.3         |   |   |   |   |
| 9  | 50004           | Minnie                    | 3.9         |   |   |   |   |
| 10 | 10008           | Pluto                     | 4           |   |   |   |   |
| 11 |                 |                           |             |   |   |   |   |
| 12 |                 |                           |             |   |   |   |   |
| 13 | <b>Courses</b>  |                           |             |   |   |   |   |
| 14 | <b>CID</b>      | <b>C-Name</b>             | <b>Room</b> |   |   |   |   |
| 15 | 1012            | CS145 - Toon systems      | Nvidia      |   |   |   |   |
| 16 | 1017            | CS161 - Animation art     | Gates 300   |   |   |   |   |
| 17 | 1019            | CS245 - Painting town red | Packard 45  |   |   |   |   |
| 18 |                 |                           |             |   |   |   |   |
| 19 |                 |                           |             |   |   |   |   |

| <b>Enrolled</b> |            |              |
|-----------------|------------|--------------|
| <b>SID</b>      | <b>CID</b> | <b>Grade</b> |
| 40001           | 1012       | A            |
| 50004           | 1012       | A-           |
| 40001           | 1017       | B+           |

## Queries ["compute" over tables]

- Minnie's GPA?
- AVG student GPA?
- Mickey's classes?
- AVG student GPA in CS145?

## Key concept

### Data model

Relational model (aka tables)

Simple and most popular

Elegant algebra (E.F. Codd et al)

Every relation has a schema

Logical Schema: describes types, names

Physical Schema: describes data layout

Virtual Schema (Views): derived tables

Data model:

Organizing principle of  
data + operations

Schema:

Describes blueprint of  
table (s)

# Data independence

1. Can we add a new column or attribute without rewriting the application?

## Logical Data Independence

Protection from changes in the Logical Structure of the data

2. Do you need to care which disks/machines are the data stored on?

## Physical Data Independence

Protection from Physical Layout Changes

Key concept

Data independence

# Python operating on Lists [reminder]

## BASIC TYPES

Int, long int  
string ...

## MAP + FILTER

`map(function, list of inputs)`

`filter(function, list of inputs)`

- Map applies function to input list
- Filter returns sub-list that satisfies a filter condition

## REDUCE/AGGREGATE

`reduce (...)`

- Reduce runs a computation on a list and returns a result. E.g., SUM, MAX, MIN

*For review, check out your favorite python tutorial (e.g, [https://book.pythontips.com/en/latest/map\\_filter.html](https://book.pythontips.com/en/latest/map_filter.html))*



# SQL Queries on Tables (Lists of rows)

## BASIC TYPES

Int32, int64  
Char[n] ...  
Float32, float64

## MAP + FILTER

Single table query

```
SELECT c1, c2  
FROM T  
WHERE condition;
```

Multi table JOIN

```
SELECT c1, c2  
FROM T1, T2  
WHERE condition;
```

## REDUCE/AGGREGATE

```
SELECT SUM(c1*c2)  
FROM T  
WHERE condition  
GROUP BY c3;
```

Map-Filter-Reduce pattern: Same simple/powerful idea in MapReduce, Hadoop, Spark, etc.

## Preview SQL queries

### QUERYING DATA FROM A TABLE

**SELECT c1, c2 FROM t;**  
Query data in columns c1, c2 from a table

**SELECT \* FROM t;**  
Query all rows and columns from a table

**SELECT c1, c2 FROM t  
WHERE condition;**  
Query data and filter rows with a condition

**SELECT DISTINCT c1 FROM t  
WHERE condition;**  
Query distinct rows from a table

**SELECT c1, c2 FROM t  
ORDER BY c1 ASC [DESC];**  
Sort the result set in ascending or descending order

**SELECT c1, c2 FROM t  
ORDER BY c1  
LIMIT n OFFSET offset;**  
Skip *offset* of rows and return the next *n* rows

**SELECT c1, aggregate(c2)  
FROM t  
GROUP BY c1;**  
Group rows using an aggregate function

**SELECT c1, aggregate(c2)  
FROM t  
GROUP BY c1  
HAVING condition;**  
Filter groups using HAVING clause

### QUERYING FROM MULTIPLE TABLES

**SELECT c1, c2  
FROM t1  
INNER JOIN t2 ON condition;**  
Inner join t1 and t2

**SELECT c1, c2  
FROM t1  
LEFT JOIN t2 ON condition;**  
Left join t1 and t2

**SELECT c1, c2  
FROM t1  
RIGHT JOIN t2 ON condition;**  
Right join t1 and t2

**SELECT c1, c2  
FROM t1  
FULL OUTER JOIN t2 ON condition;**  
Perform full outer join

**SELECT c1, c2  
FROM t1  
CROSS JOIN t2;**  
Produce a Cartesian product of rows in tables

**SELECT c1, c2  
FROM t1, t2;**  
Another way to perform cross join

**SELECT c1, c2  
FROM t1 A  
INNER JOIN t2 B ON condition;**  
Join t1 to itself using INNER JOIN clause

### USING SQL OPERATORS

**SELECT c1, c2 FROM t1  
UNION [ALL]  
SELECT c1, c2 FROM t2;**  
Combine rows from two queries

**SELECT c1, c2 FROM t1  
INTERSECT  
SELECT c1, c2 FROM t2;**  
Return the intersection of two queries

**SELECT c1, c2 FROM t1  
MINUS  
SELECT c1, c2 FROM t2;**  
Subtract a result set from another result set

**SELECT c1, c2 FROM t1  
WHERE c1 [NOT] LIKE pattern;**  
Query rows using pattern matching %, \_

**SELECT c1, c2 FROM t  
WHERE c1 [NOT] IN value\_list;**  
Query rows in a list

**SELECT c1, c2 FROM t  
WHERE c1 BETWEEN low AND high;**  
Query rows between two values

**SELECT c1, c2 FROM t  
WHERE c1 IS [NOT] NULL;**  
Check if values in a table is NULL or not



# Relational Algebra

# Operations in Relational Algebra

- C1- The usual set operations - union, intersection, and difference – applied to relations.
- C2 - Operations that remove parts of a relation: "selection" eliminates some rows (tuples), and "projection" eliminates some columns.
- C3 - Operations that combine the tuples of two relations, including "Cartesian product," which pairs the tuples of two relations in all possible ways, and various kinds of "join" operations, which selectively pair tuples from two relations.
- C4 - An operation called "renaming" that does not affect the tuples of a relation, but changes the relation schema, i.e., the names of the attributes and/or the name of the relation itself.



# Single - table queries

## Relational Algebra(C1)

# Selection

- Returns all tuples which satisfy a condition
- Notation:  $\sigma_c(R)$
- Examples
  - $\sigma_{\text{Salary} > 40000}(\text{Employee})$
  - $\sigma_{\text{name} = \text{"Smith"}}(\text{Employee})$
- The condition  $c$  can be  $=, <, \leq, >, \geq, <>$

| SSN     | Name  | Salary |
|---------|-------|--------|
| 1234545 | John  | 200000 |
| 5423341 | Smith | 600000 |
| 4352342 | Fred  | 500000 |

$\sigma_{\text{Salary} > 40000}$  (Employee)

| SSN     | Name  | Salary |
|---------|-------|--------|
| 5423341 | Smith | 600000 |
| 4352342 | Fred  | 500000 |

# Projection

- Eliminates columns, then removes duplicates
- Notation:  $\Pi_{A1, \dots, An}(R)$
- Example: project social-security number and names:
  - $\Pi_{SSN, Name}(\text{Employee})$
  - Output schema: Answer(SSN, Name)



| SSN     | Name | Salary |
|---------|------|--------|
| 1234545 | John | 200000 |
| 5423341 | John | 600000 |
| 4352342 | John | 200000 |

$\Pi_{\text{Name,Salary}}$  (Employee)

| Name | Salary |
|------|--------|
| John | 20000  |
| John | 60000  |

What you will  
learn about  
in this section

1. The SFW query
2. Other useful operators: LIKE, DISTINCT, ORDER BY



# SQL Query

Basic form (there are many many more bells and whistles)

**SELECT** <attributes>

**FROM** <one or more relations>

**WHERE** <conditions>

Call this a **SFW** query.

# Simple SQL Query: Selection

**Selection** is the operation of filtering a relation's tuples on some condition

```
SELECT *  
FROM Product  
WHERE Category = 'Gadgets'
```

| PName      | Price   | Category | Manuf  |
|------------|---------|----------|--------|
| Gizmo      | \$19.99 | Gadgets  | GWorks |
| Powergizmo | \$29.99 | Gadgets  | GWorks |

|             |          |             |         |
|-------------|----------|-------------|---------|
| SingleTouch | \$149.99 | Photography | Canon   |
| MultiTouch  | \$203.99 | Household   | Hitachi |



| PName      | Price   | Category | Manuf  |
|------------|---------|----------|--------|
| Gizmo      | \$19.99 | Gadgets  | GWorks |
| Powergizmo | \$29.99 | Gadgets  | GWorks |

# Simple SQL Query: Projection

**Projection** is the operation of producing an output table with tuples that have a subset of their prior attributes

```
SELECT Pname, Price,  
Manufacturer  
FROM Product  
WHERE Category = 'Gadgets'
```

| PName      | Price   | Category | Manuf  |
|------------|---------|----------|--------|
| Gizmo      | \$19.99 | Gadgets  | GWorks |
| Powergizmo | \$29.99 | Gadgets  | GWorks |

|             |          |             |         |
|-------------|----------|-------------|---------|
| SingleTouch | \$149.99 | Photography | Canon   |
| MultiTouch  | \$203.99 | Household   | Hitachi |



| PName      | Price   | Manuf  |
|------------|---------|--------|
| Gizmo      | \$19.99 | GWorks |
| Powergizmo | \$29.99 | GWorks |

# Notation

Input Schema

Product(PName, Price, Category,  
Manufacturer)

```
SELECT Pname, Price,  
       Manufacturer  
FROM   Product  
WHERE  Category = 'Gadgets'
```

Output Schema

Answer(PName, Price, Manufacturer)







# A Few Details

- SQL **commands** are case insensitive:  
Same: SELECT, Select, select  
Same: Product, product
- **Values are not:**  
Different: 'Seattle', 'seattle'
- Use single quotes for constants:  
'abc' - yes  
"abc" - no



## LIKE: Simple String Pattern Matching

```
SELECT *  
FROM Products  
WHERE PName LIKE '%gizmo%'
```

- s **LIKE** p: pattern matching on strings
- p may contain two special symbols:
  - % = any sequence of characters
  - \_ = any single character



# DISTINCT: Eliminating Duplicates

```
SELECT DISTINCT Category  
FROM Product
```



Category


Gadgets

Photography

Household

**Versus**

```
SELECT Category  
FROM Product
```



Category

Gadgets

Gadgets

Photography

Household



# ORDER BY: Sorting the Results

```
SELECT PName, Price, Manufacturer
FROM Product
WHERE Category='gizmo' AND Price > 50
ORDER BY Price, PName
```

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.



# SQL Part II

## JOIN, Aggregate Functions, Group By and HAVING

## Preview SQL queries

### QUERYING DATA FROM A TABLE

**SELECT c1, c2 FROM t;**  
Query data in columns c1, c2 from a table

**SELECT \* FROM t;**  
Query all rows and columns from a table

**SELECT c1, c2 FROM t  
WHERE condition;**  
Query data and filter rows with a condition

**SELECT DISTINCT c1 FROM t  
WHERE condition;**  
Query distinct rows from a table

**SELECT c1, c2 FROM t  
ORDER BY c1 ASC [DESC];**  
Sort the result set in ascending or descending order

**SELECT c1, c2 FROM t  
ORDER BY c1  
LIMIT n OFFSET offset;**  
Skip *offset* of rows and return the next *n* rows

**SELECT c1, aggregate(c2)  
FROM t  
GROUP BY c1;**  
Group rows using an aggregate function

**SELECT c1, aggregate(c2)  
FROM t  
GROUP BY c1  
HAVING condition;**  
Filter groups using HAVING clause

### QUERYING FROM MULTIPLE TABLES

**SELECT c1, c2  
FROM t1  
INNER JOIN t2 ON condition;**  
Inner join t1 and t2

**SELECT c1, c2  
FROM t1  
LEFT JOIN t2 ON condition;**  
Left join t1 and t2

**SELECT c1, c2  
FROM t1  
RIGHT JOIN t2 ON condition;**  
Right join t1 and t2

**SELECT c1, c2  
FROM t1  
FULL OUTER JOIN t2 ON condition;**  
Perform full outer join

**SELECT c1, c2  
FROM t1  
CROSS JOIN t2;**  
Produce a Cartesian product of rows in tables

**SELECT c1, c2  
FROM t1, t2;**  
Another way to perform cross join

**SELECT c1, c2  
FROM t1 A  
INNER JOIN t2 B ON condition;**  
Join t1 to itself using INNER JOIN clause

### USING SQL OPERATORS

**SELECT c1, c2 FROM t1  
UNION [ALL]  
SELECT c1, c2 FROM t2;**  
Combine rows from two queries

**SELECT c1, c2 FROM t1  
INTERSECT  
SELECT c1, c2 FROM t2;**  
Return the intersection of two queries

**SELECT c1, c2 FROM t1  
MINUS  
SELECT c1, c2 FROM t2;**  
Subtract a result set from another result set

**SELECT c1, c2 FROM t1  
WHERE c1 [NOT] LIKE pattern;**  
Query rows using pattern matching %, \_


**SELECT c1, c2 FROM t  
WHERE c1 [NOT] IN value\_list;**  
Query rows in a list

**SELECT c1, c2 FROM t  
WHERE c1 BETWEEN low AND high;**  
Query rows between two values

**SELECT c1, c2 FROM t  
WHERE c1 IS [NOT] NULL;**  
Check if values in a table is NULL or not



# Reminder on schemas



Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Students(sid: string, name: string, gpa: float)

Enrolled(student\_id: string, cid: string, grade: string)

We'll use different  
Tables/tuples, for examples  
to build ideas

Data about local areas (for real-world examples)

SolarPanel(region\_name: string, kw\_total: float, carbon\_offset\_ton\_metrics: float, ...)

Census(zipcode: string, population: int, ...)

Pollution(zipcode: string, Particle\_count: int...)

BikeShare(zipcode: string, trip\_origin: float, trip\_end: float, ...)

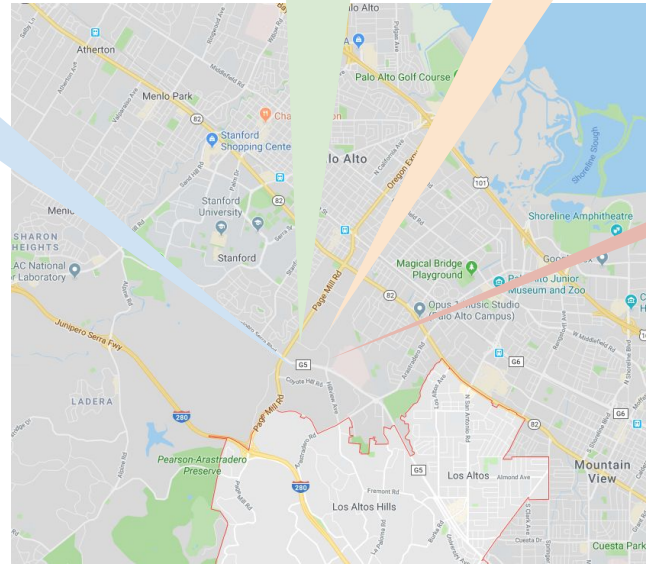
...



## A detailed image of a spiral galaxy, likely the Milky Way, showing its central bulge and prominent spiral arms. The galaxy is oriented diagonally, with the bright central core at the top left. The spiral arms are composed of numerous stars and interstellar dust, appearing as a complex, swirling pattern of light and dark regions. The background is a deep black, dotted with distant stars and faint nebulae.

| Census  |                   |
|---------|-------------------|
| Zipcode | Population Census |
| 94305   | 14301             |
| 94040   | 20301             |
| 94041   | 189               |

| Pollution |                |
|-----------|----------------|
| Zipcode   | Particle count |
| 94305     | 40             |
| 94040     | 22             |
| 94041     | 57             |



| OmniData   |          |               |     |            |                |                      |        |
|------------|----------|---------------|-----|------------|----------------|----------------------|--------|
| SolarPanel |          |               |     | Census     | Pollution      | Bike share locations |        |
| Zipcode    | KW-Total | Carbon offset | ... | Population | Particle count | Lat                  | Lng    |
| 94305      | 14.4     | 29            |     | 14301      | 40             | 35.1                 | 122.12 |
| 94305      |          |               |     |            |                | 35.2                 | 122.13 |
| 94305      |          |               |     |            |                | 35.2                 | 122.1  |
| 94305      |          |               |     |            |                | 35.1                 | 122.12 |
| 94305      |          |               |     |            |                | ...                  | ...    |

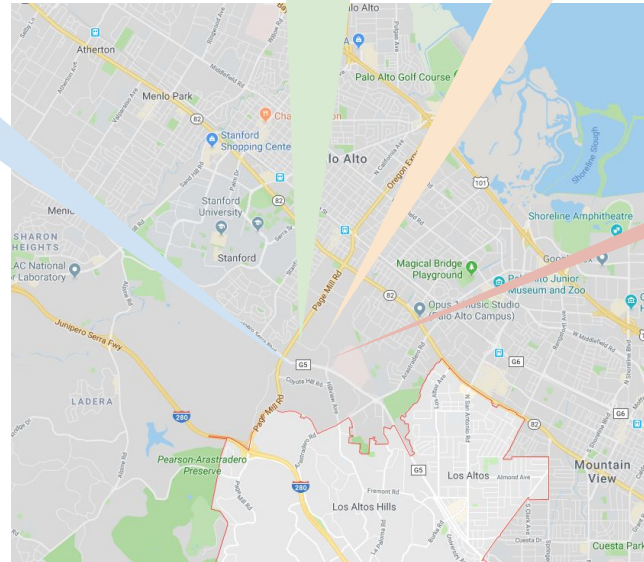
## Option 1: A few “good” tables (with 10s of columns)

| SolarPanel |          |               |     |
|------------|----------|---------------|-----|
| Zipcode    | KW-Total | Carbon offset | ... |
| 94305      | 14.4     | 29            |     |
| 94040      | 32.1     | 42            |     |
| 94041      | 29.1     | 37.38         |     |

| Census  |                   |  |
|---------|-------------------|--|
| Zipcode | Population Census |  |
| 94305   | 14301             |  |
| 94040   | 20301             |  |
| 94041   | 189               |  |

| Bike share locations |      |        |  |
|----------------------|------|--------|--|
| Zipcode              | Lat  | Lng    |  |
| 94305                | 35.1 | 122.12 |  |
| 94305                | 35.2 | 122.13 |  |
| 94041                | 35.1 | 121.27 |  |
| 94041                |      |        |  |
| 94041                |      |        |  |

| Pollution |                |
|-----------|----------------|
| Zipcode   | Particle count |
| 94305     | 40             |
| 94040     | 22             |
| 94041     | 57             |



Trade offs?

- Reads? Writes?
- 100s - thousands of applications reading/writing data

⇒ Hybrids: 1 column → all columns

(Week 7: What's a good schema design?)

## Option 2: ‘FrankenTable’ (with 1000s of columns)

| Omnidata   |          |               |     |            |  |                |      |
|------------|----------|---------------|-----|------------|--|----------------|------|
| SolarPanel |          |               |     | Census     |  | Pollution      |      |
| Zipcode    | KW-Total | Carbon offset | ... | Population |  | Particle count | Lat  |
| 94305      | 14.4     | 29            |     | 14301      |  | 40             | 35.1 |
| 94305      | 14.4     | 29            |     | 14301      |  | 40             | 35.2 |
| 94305      | 14.4     | 29            |     | 14301      |  | 40             | 35.1 |
| 94305      | 14.4     | 29            |     | 14301      |  | 40             | 35.1 |

# Assume

(for now)

- ▶ Assume we have a set of “good” tables
  - ▷ How do we “connect” (join) those tables?
- ▶ Related important question
  - ▷ How to break up a “Franken” Table into “good” Tables? (i.e, design “good” schema)
  - ▷ Study in Week 5,6,7





## Basic Example of Join

# Joins

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

*Ex:* Find all products under \$200  
manufactured in Japan; return their names  
and prices.

Note: we will often omit  
attribute types in  
schema definitions for  
brevity, but assume  
attributes are always  
atomic types

# Joins

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Ex: Find all products under \$200  
manufactured in Japan; return their names  
and prices.

SELECT PName, Price

FROM Product, Company

WHERE Manufacturer = CName  
AND Country='Japan'  
AND Price <= 200

PName

Price

Category

Manufacturer

Gizmo

\$19

Gadgets

GizmoWorks

Powergizmo

\$29

Gadgets

GizmoWorks

SingleTouch

\$149

Photography

Canon

MultiTouch

\$203

Household

Hitachi

A join between  
tables returns all  
unique  
combinations of  
their tuples **which**  
**meet some**  
**specified join**  
**condition**

CName

Stock  
Price

Countr  
y

GizmoWorks

25

USA

Canon

65

Japan

Hitachi

15

Japan

# Joins

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Several equivalent ways to write a basic join  
in SQL:

```
SELECT PName, Price
FROM   Product, Company
WHERE  Manufacturer =
      CName
      AND Country='Japan'
      AND Price <= 200
```

```
SELECT PName, Price
FROM   Product
JOIN   Company
ON     Manufacturer =
      CName
WHERE  Price <= 200
      AND Country='Japan'
```

A few more later on

# Joins

## Product

| <u>PName</u> | Price | Category    | Manufacturer |
|--------------|-------|-------------|--------------|
| Gizmo        | \$19  | Gadgets     | GizmoWorks   |
| Powergizmo   | \$29  | Gadgets     | GizmoWorks   |
| SingleTouch  | \$149 | Photography | Canon        |
| MultiTouch   | \$205 | Household   | Hitachi      |

## Company

| <u>CName</u> | Stock Price | Country |
|--------------|-------------|---------|
| GizmoWorks   | 25          | USA     |
| Canon        | 65          | Japan   |
| Hitachi      | 15          | Japan   |

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

**PName      Price**

**SingleTouch      \$149**

# Tuple Variable Ambiguity in Multi-Table

Person(name, address, worksfor)  
Company(name, address)

1. **SELECT DISTINCT** name, address
2. **FROM** Person, Company
3. **WHERE** worksfor = name

Which “address”  
does this refer to?

**Which name”s??**

# Tuple Variable Ambiguity in Multi-Table

Both  
equivalent  
ways to  
resolve  
variable  
ambiguity

```
Person(name, address, worksfor)  
Company(name, address)
```

```
SELECT DISTINCT Person.name,  
Person.address  
FROM      Person, Company  
WHERE      Person.worksfor =  
Company.name
```

```
SELECT DISTINCT p.name, p.address  
FROM      Person p, Company c  
WHERE      p.worksfor = c.name
```

# Semantics of JOINS

```
SELECT  $x_1.a_1, x_1.a_2, \dots, x_n.a_k$   
FROM  $R_1$  AS  $x_1, R_2$  AS  $x_2, \dots,$   
 $R_n$  AS  $x_n$   
WHERE  $\text{Conditions}(x_1, \dots, x_n)$ 
```

Answer = { }

```
for  $x_1$  in  $R_1$  do  
  for  $x_2$  in  $R_2$  do  
    .....  
    for  $x_n$  in  $R_n$  do  
      if  $\text{Conditions}(x_1, \dots, x_n)$   
        then Answer = Answer  $\cup \{(x_1.a_1, x_1.a_2, \dots,$   
 $x_n.a_k)\}$   
return Answer
```

**Note:**  
This is a  
*multiset*  
union



```
SELECT R.A  
FROM R, S  
WHERE R.A = S.B
```

- Take **cross product**

$$V = R \times S$$

- Apply **selections/conditions**

$$Y = \{(r, s) \text{ in } V \mid r.A == s.B\}$$

- Apply **projections** to get final output

$$Z = (y.A) \text{ for } y \text{ in } Y$$

Recall: Cross product ( $R \times S$ ) is the set of all unique tuples in  $R, S$

Ex:  $\{a, b, c\} \times \{1, 2\}$

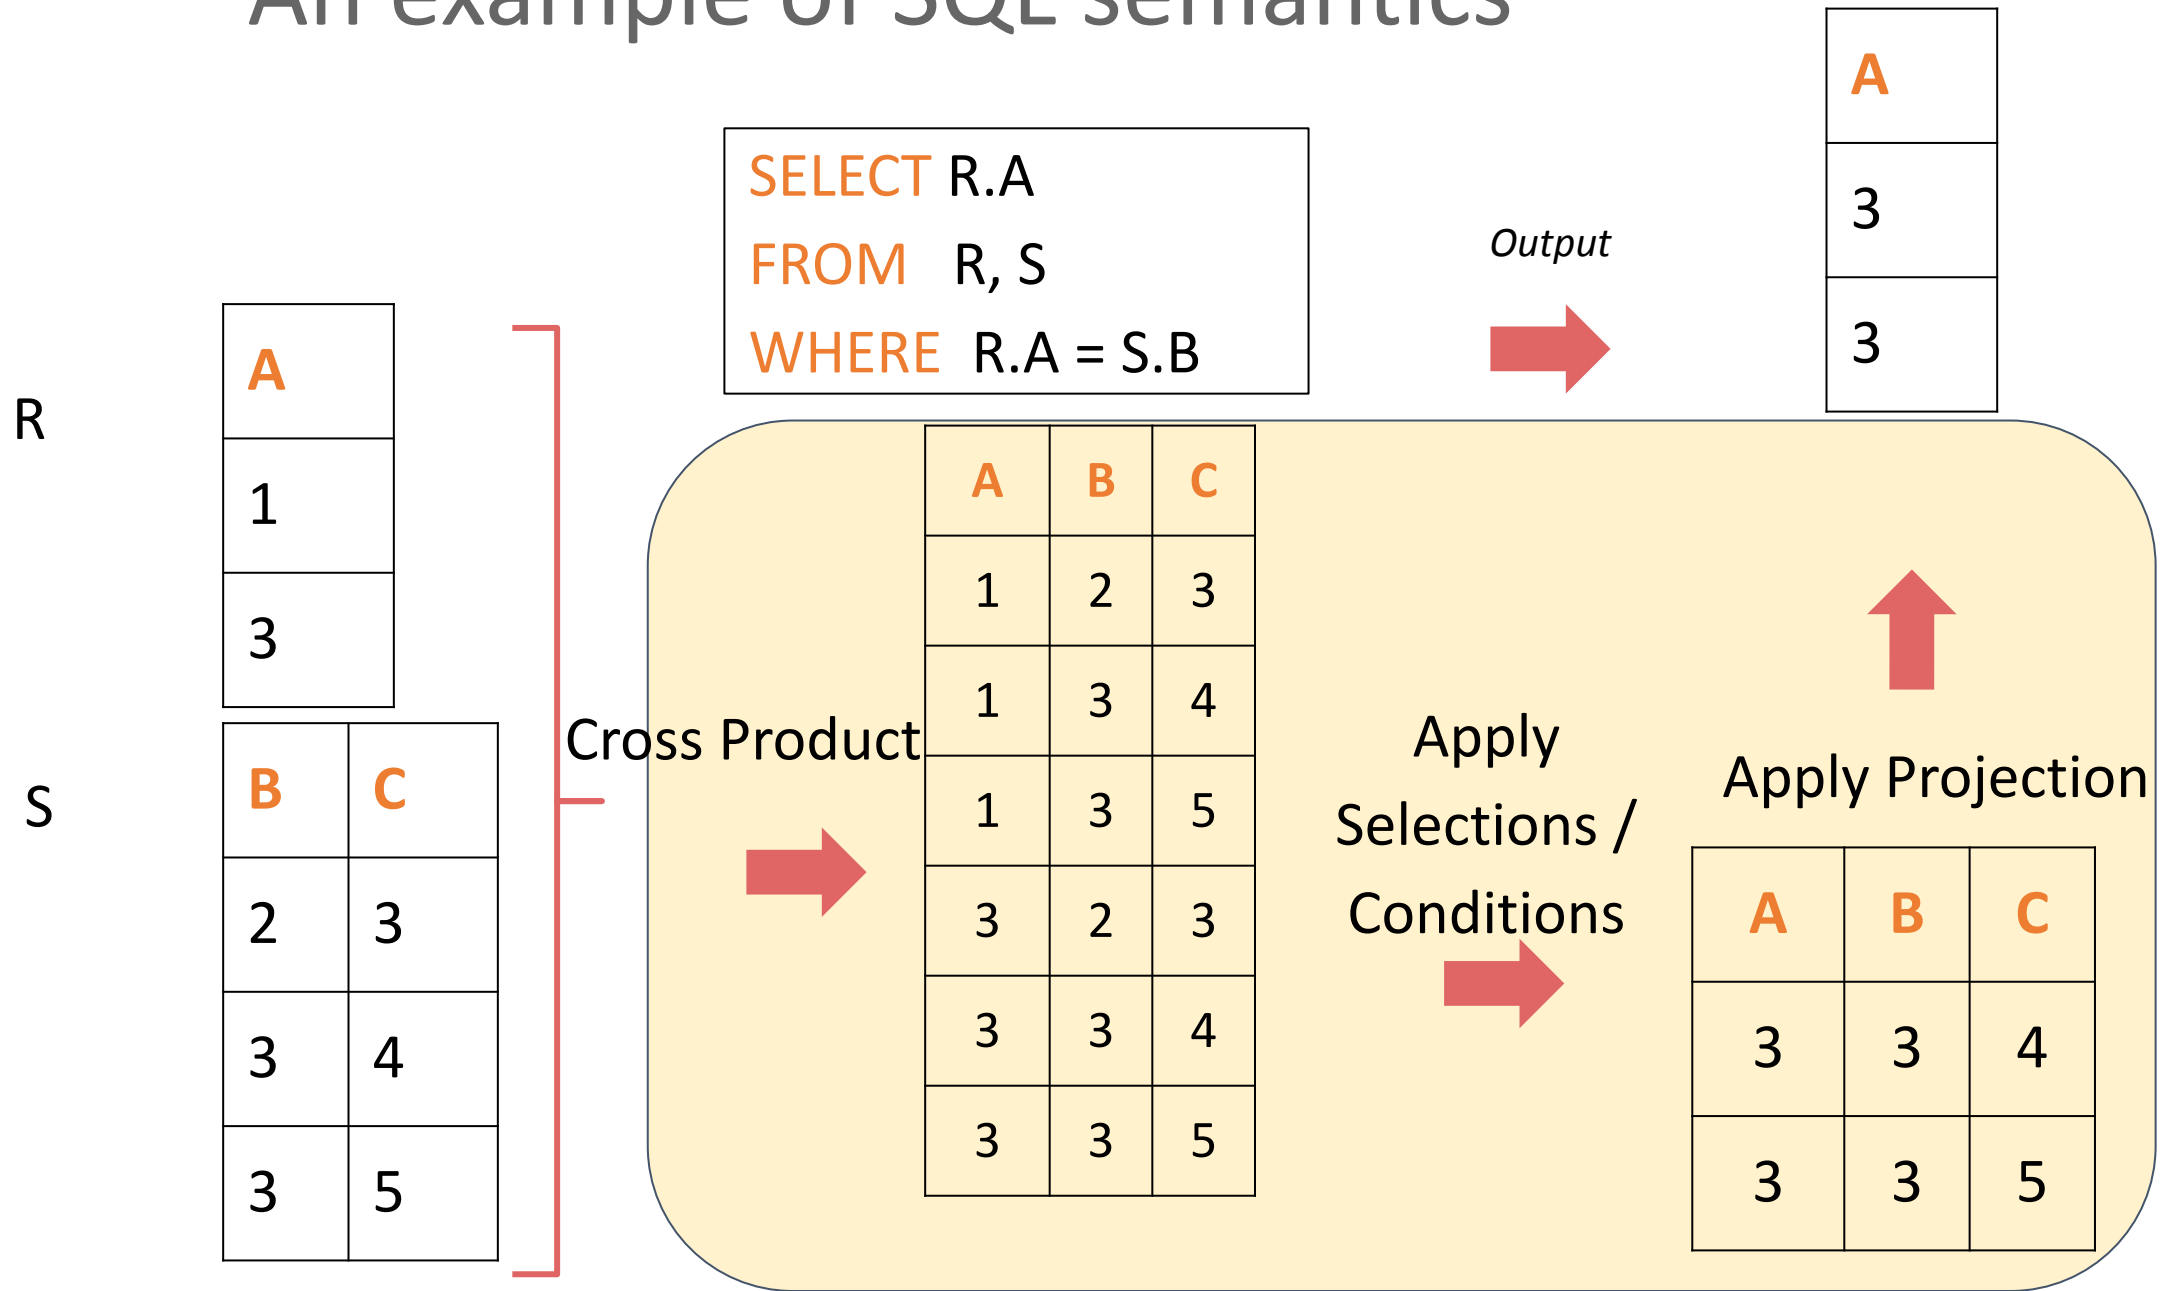
$= \{(a, 1), (a, 2), (b, 1), (b, 2), (c, 1), (c, 2)\}$

= Filtering!

= Returning only *some* attributes

Remembering this order is critical to understanding the output of certain queries (see later on...)

# An example of SQL semantics



Note: we say “semantics” not “execution order”

- The preceding slides show *what a join means*
- Not actually how the DBMS executes it under the covers

# A Subtlety about Joins

```
Product(PName, Price, Category,  
Manufacturer)  
Company(CName, StockPrice, Country)
```

Find all countries that manufacture some product in the 'Gadgets' category.

```
SELECT Country  
FROM Product, Company  
WHERE Manufacturer=CName AND  
Category='Gadgets'
```

# A Subtlety about Joins

## Product

| PName       | Price | Category    | Manuf   |
|-------------|-------|-------------|---------|
| Gizmo       | \$19  | Gadgets     | GWorks  |
| Powergizmo  | \$29  | Gadgets     | GWorks  |
| SingleTouch | \$149 | Photography | Canon   |
| MultiTouch  | \$203 | Household   | Hitachi |

## Company

| Cname   | Stock | Country |
|---------|-------|---------|
| GWorks  | 25    | USA     |
| Canon   | 65    | Japan   |
| Hitachi | 15    | Japan   |



```
SELECT Country
FROM Product, Company
WHERE Manufacturer=Cname
AND Category='Gadgets'
```


Country

USA<sup>?</sup>

?  
USA

Ans? DISTINCT

What is the Problem? What is the Solution?





# JOIN TYPES and NULL OPERATION

Inner, Full, Left and Right Outer Join

# NULLS in SQL

- Whenever we don't have a value, we can put a NULL
- Can mean many things:
  - Value does not exist
  - Value exists but is unknown
  - Value not applicable
  - Etc.
- The schema specifies for each attribute if can be null (*nullable* attribute) or not
- How does SQL cope with tables that have NULLs?



# Null Values

- *For numerical operations, NULL -> NULL:*
  - If  $x = \text{NULL}$  then  $4 \cdot (3 - x) / 7$  is still NULL
- *For boolean operations, in SQL there are three values:*

|                |          |            |
|----------------|----------|------------|
| <b>FALSE</b>   | <b>=</b> | <b>0</b>   |
| <b>UNKNOWN</b> | <b>=</b> | <b>0.5</b> |
| <b>TRUE</b>    | <b>=</b> | <b>1</b>   |

- If  $x = \text{NULL}$  then  $x == \text{"Joe"}$  is UNKNOWN (Is  $x$  equal to 'Joe'?)

# Null Values

- $C1 \text{ AND } C2 = \min(C1, C2)$
- $C1 \text{ OR } C2 = \max(C1, C2)$
- $\text{NOT } C1 = 1 - C1$

```
SELECT *  
FROM Person  
WHERE (age < 25)  
      AND (height > 6 AND weight > 190)
```

Won't return e.g.  
(age=20  
height=NULL  
weight=200)!

Rule in SQL: include only tuples that yield TRUE (1.0)

# Null Values

Unexpected behavior:

```
SELECT *  
FROM Person  
WHERE age < 25 OR age >= 25
```

Some Persons are not included !

# Null Values

Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

```
SELECT *  
FROM Person  
WHERE age < 25 OR age >= 25  
      OR age IS NULL
```

Now it includes all Persons!

# RECAP: Joins


By default, joins in SQL are “**inner joins**”:

```
Product(name, category)
Purchase(prodName, store)
```

```
SELECT Product.name, Purchase.store
FROM   Product, Purchase
WHERE  Product.name = Purchase.prodName
```

```
SELECT Product.name, Purchase.store
FROM   Product
JOIN   Purchase ON Product.name = Purchase.prodName
```

```
SELECT Product.name, Purchase.store
FROM   Product
INNER JOIN Purchase
ON     Product.name = Purchase.prodName
```



Both equivalent:  
Both INNER JOINS!

[Like Below]

# Inner Joins + NULLS = Lost data?


By default, joins in SQL are “**inner joins**”:

```
Product(name, category)
Purchase(prodName, store)
```

```
SELECT Product.name, Purchase.store
FROM Product
JOIN Purchase ON Product.name = Purchase.prodName
```

| Product    |          |
|------------|----------|
| Name       | Category |
| Iphone     | Media    |
| Roomba     | Cleaner  |
| Ford Pinto | Car      |
| Tesla      | Car      |

| Purchase |             |
|----------|-------------|
| ProdName | Store       |
| Iphone   | Apple Store |
| Tesla    | Tesla Store |



# Outer Joins

- An **outer join** returns tuples from the joined relations that don't have a corresponding tuple in the other relations
  - I.e. If we join relations A and B on  $a.X = b.X$ , and there is an entry in A with  $X=5$ , but none in B with  $X=5$ ...
  - A LEFT OUTER JOIN will return a tuple (a, NULL)!
- Left outer joins in SQL:

```
SELECT Product.name, Purchase.store  
FROM   Product  
LEFT OUTER JOIN Purchase ON  
        Product.name = Purchase.prodName
```

Now we'll get products even if they didn't sell

# LEFT OUTER JOIN

**Product**

| name       | category |
|------------|----------|
| iphone     | media    |
| Tesla      | car      |
| Ford Pinto | car      |

**Purchase**

| prodName | store       |
|----------|-------------|
| iPhone   | Apple store |
| Tesla    | car         |
| iPhone   | Apple store |

```
SELECT Product.name, Purchase.store
FROM Product
LEFT OUTER JOIN Purchase
ON Product.name = Purchase.prodName
```



| name       | store       |
|------------|-------------|
| iPhone     | Apple store |
| iPhone     | Apple store |
| Tesla      | car         |
| Ford Pinto | NULL        |



# Other Outer Joins

- Left outer join:
  - Include the left tuple even if there's no match
- Right outer join:
  - Include the right tuple even if there's no match
- Full outer join:
  - Include the both left and right tuples even if there's no match

# Natural Join

| <i>A</i> | <i>B</i> |
|----------|----------|
| 1        | 2        |
| 3        | 4        |

(a) Relation *R*

| <i>A</i> | <i>B</i> | <i>C</i> | <i>D</i> |
|----------|----------|----------|----------|
| 1        | 2        | 5        | 6        |
| 3        | 4        | 7        | 8        |

| <i>B</i> | <i>C</i> | <i>D</i> |
|----------|----------|----------|
| 2        | 5        | 6        |
| 4        | 7        | 8        |
| 9        | 10       | 11       |

(b) Relation *S*

| <i>A</i> | <i>R.B</i> | <i>S.B</i> | <i>C</i> | <i>D</i> |
|----------|------------|------------|----------|----------|
| 1        | 2          | 2          | 5        | 6        |
| 1        | 2          | 4          | 7        | 8        |
| 1        | 2          | 9          | 10       | 11       |
| 3        | 4          | 2          | 5        | 6        |
| 3        | 4          | 4          | 7        | 8        |
| 3        | 4          | 9          | 10       | 11       |

(c) Result  $R \times S$

| SR.NO. | NATURAL JOIN   | INNER JOIN  |
|--------|--|---|
| 1.     | Natural Join joins two tables based on same attribute name and datatypes.  | Inner Join joins two table on the basis of the column which is explicitly specified in the ON clause.                 |
| 2.     | In Natural Join, The resulting table will contain all the attributes of both the tables but keep only one copy of each common column | In Inner Join, The resulting table will contain all the attribute of both the tables including duplicate columns also |
| 3.     | In Natural Join, If there is no condition specifies then it returns the rows based on the common column                              | In Inner Join, only those records will return which exists in both the tables   |
| 4.     | SYNTAX:<br>SELECT *<br>FROM table1 NATURAL JOIN table2;  | SYNTAX:<br>SELECT *<br>FROM table1 INNER JOIN table2 ON<br>table1.Column_Name =<br>table2.Column_Name;                |

## Difference between INNER and NATURAL JOIN

# Theta Join

| <i>A</i> | <i>B</i> | <i>C</i> |
|----------|----------|----------|
| 1        | 2        | 3        |
| 6        | 7        | 8        |
| 9        | 7        | 8        |

(a) Relation *U*

| <i>B</i> | <i>C</i> | <i>D</i> |
|----------|----------|----------|
| 2        | 3        | 4        |
| 2        | 3        | 5        |
| 7        | 8        | 10       |

(b) Relation *V*

| <i>A</i> | <i>B</i> | <i>C</i> | <i>D</i> |
|----------|----------|----------|----------|
| 1        | 2        | 3        | 4        |
| 1        | 2        | 3        | 5        |
| 6        | 7        | 8        | 10       |
| 9        | 7        | 8        | 10       |

(c) Result  $U \bowtie V$

$$U \bowtie_{A < D} V$$

| <i>A</i> | <i>U.B</i> | <i>U.C</i> | <i>V.B</i> | <i>V.C</i> | <i>D</i> |
|----------|------------|------------|------------|------------|----------|
| 1        | 2          | 3          | 2          | 3          | 4        |
| 1        | 2          | 3          | 2          | 3          | 5        |
| 1        | 2          | 3          | 7          | 8          | 10       |
| 6        | 7          | 8          | 7          | 8          | 10       |
| 9        | 7          | 8          | 7          | 8          | 10       |

$$U \bowtie_{A < D \text{ AND } U.B \neq V.B} V$$

| <i>A</i> | <i>U.B</i> | <i>U.C</i> | <i>V.B</i> | <i>V.C</i> | <i>D</i> |
|----------|------------|------------|------------|------------|----------|
| 1        | 2          | 3          | 7          | 8          | 10       |

# Self Join

```
SELECT column_name(s)
FROM table1 T1, table1 T2
WHERE condition;
```

```
SELECT A.CustomerName AS CustomerName1, B.CustomerName AS CustomerName2, A.City
FROM Customers A, Customers B
WHERE A.CustomerID <> B.CustomerID
AND A.City = B.City
ORDER BY A.City;
```

w3schools

| CustomerID | CustomerName                       | ContactName    | Address                       | City        | PostalCode | Country |
|------------|------------------------------------|----------------|-------------------------------|-------------|------------|---------|
| 1          | Alfreds Futterkiste                | Maria Anders   | Obere Str. 57                 | Berlin      | 12209      | Germany |
| 2          | Ana Trujillo Emparedados y helados | Ana Trujillo   | Avda. de la Constitución 2222 | México D.F. | 05021      | Mexico  |
| 3          | Antonio Moreno Taquería            | Antonio Moreno | Mataderos 2312                | México D.F. | 05023      | Mexico  |

| CustomerName1              | CustomerName2              | City         |
|----------------------------|----------------------------|--------------|
| Cactus Comidas para llevar | Océano Atlántico Ltda.     | Buenos Aires |
| Cactus Comidas para llevar | Rancho grande              | Buenos Aires |
| Océano Atlántico Ltda.     | Cactus Comidas para llevar | Buenos Aires |
| Océano Atlántico Ltda.     | Rancho grande              | Buenos Aires |
| Rancho grande              | Cactus Comidas para llevar | Buenos Aires |

# EquiJoin and Natural Join

- Equi Join is a join using one common column (referred to in the “on” clause)
- Natural Join is an implicit join clause based on the common columns in the two tables being joined. Common columns are columns that have the same name in both tables.

## Preview SQL queries

### QUERYING DATA FROM A TABLE

**SELECT c1, c2 FROM t;**  
Query data in columns c1, c2 from a table

**SELECT \* FROM t;**  
Query all rows and columns from a table

**SELECT c1, c2 FROM t  
WHERE condition;**  
Query data and filter rows with a condition

**SELECT DISTINCT c1 FROM t  
WHERE condition;**  
Query distinct rows from a table

**SELECT c1, c2 FROM t  
ORDER BY c1 ASC [DESC];**  
Sort the result set in ascending or descending order

**SELECT c1, c2 FROM t  
ORDER BY c1  
LIMIT n OFFSET offset;**  
Skip *offset* of rows and return the next *n* rows

**SELECT c1, aggregate(c2)  
FROM t  
GROUP BY c1;**  
Group rows using an aggregate function

**SELECT c1, aggregate(c2)  
FROM t  
GROUP BY c1  
HAVING condition;**  
Filter groups using HAVING clause

### QUERYING FROM MULTIPLE TABLES

**SELECT c1, c2  
FROM t1  
INNER JOIN t2 ON condition;**  
Inner join t1 and t2

**SELECT c1, c2  
FROM t1  
LEFT JOIN t2 ON condition;**  
Left join t1 and t2

**SELECT c1, c2  
FROM t1  
RIGHT JOIN t2 ON condition;**  
Right join t1 and t2

**SELECT c1, c2  
FROM t1  
FULL OUTER JOIN t2 ON condition;**  
Perform full outer join

**SELECT c1, c2  
FROM t1  
CROSS JOIN t2;**  
Produce a Cartesian product of rows in tables

**SELECT c1, c2  
FROM t1, t2;**  
Another way to perform cross join

**SELECT c1, c2  
FROM t1 A  
INNER JOIN t2 B ON condition;**  
Join t1 to itself using INNER JOIN clause

### USING SQL OPERATORS

**SELECT c1, c2 FROM t1  
UNION [ALL]  
SELECT c1, c2 FROM t2;**  
Combine rows from two queries

**SELECT c1, c2 FROM t1  
INTERSECT  
SELECT c1, c2 FROM t2;**  
Return the intersection of two queries

**SELECT c1, c2 FROM t1  
MINUS  
SELECT c1, c2 FROM t2;**  
Subtract a result set from another result set

**SELECT c1, c2 FROM t1  
WHERE c1 [NOT] LIKE pattern;**  
Query rows using pattern matching %, \_

**SELECT c1, c2 FROM t  
WHERE c1 [NOT] IN value\_list;**  
Query rows in a list

**SELECT c1, c2 FROM t  
WHERE c1 BETWEEN low AND high;**  
Query rows between two values

**SELECT c1, c2 FROM t  
WHERE c1 IS [NOT] NULL;**  
Check if values in a table is NULL or not

# Aggregation

```
SELECT AVG(price)
FROM Product
WHERE maker = "Toyota"
```

```
SELECT COUNT(*)
FROM Product
WHERE year > 1995
```

- SQL supports several **aggregation** operations:
  - SUM, COUNT, MIN, MAX, AVG



# Aggregation: COUNT

- COUNT applies to duplicates, unless otherwise stated

```
SELECT COUNT(category)
FROM Product
WHERE year > 1995
```

We probably want:

```
SELECT COUNT(DISTINCT category)
FROM Product
WHERE year > 1995
```


# Simple Aggregations

Purchase

| Product | Date  | Price | Quantity |
|---------|-------|-------|----------|
| bagel   | 10/21 | 1     | 20       |
| banana  | 10/3  | 0.5   | 10       |
| banana  | 10/10 | 1     | 10       |
| bagel   | 10/25 | 1.50  | 20       |

Example 1


```
SELECT SUM(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```



50  
(= 1\*20 + 1.50\*20)

Example 2

```
SELECT SUM(price * quantity)
FROM Purchase
```



65  
(= 1\*20 + 1.50\*20 + 0.5\*10 + 1\*10)

# Grouping and Aggregation



What GROUPings are possible?

- Type, Size, Color
- Number of holes
- Combination?

# What GROUPings are possible?

## Purchase

| Product | Date  | Price | Quantity |
|---------|-------|-------|----------|
| bagel   | 10/21 | 1     | 20       |
| banana  | 10/3  | 0.5   | 10       |
| banana  | 10/10 | 1     | 10       |
| bagel   | 10/25 | 1.50  | 20       |

## Possible Groups

- Product? (e.g. SUM(quantity) by product) # product units sold
- Date? (e.g., SUM(price\*quantity) by date) # daily sales
- Price?
- Product, Date?

# Grouping and Aggregation

```
Purchase(product, date, price, quantity)
```

```
SELECT product,  
        SUM(price * quantity) AS TotalSales  
FROM    Purchase  
WHERE   date > '10/1/2005'  
GROUP BY product
```

Find total sales  
after 10/1/2005  
per product.

Let's see what this means...

# Grouping and Aggregation

```
SELECT product,  
        SUM(price * quantity) AS TotalSales  
FROM Purchase  
WHERE date > '10/1/2005'  
GROUP BY product
```


## Semantics of the query:

1. Compute the **FROM** and **WHERE** clauses
2. Group by the attributes in the **GROUP BY**
3. Compute the **SELECT** clause: grouped attributes and aggregates

# 1. Compute the **FROM** and **WHERE** clauses

```
SELECT product, SUM(price*quantity) AS TotalSales  
FROM Purchase  
WHERE date > '10/1/2005'  
GROUP BY product
```

**FROM**



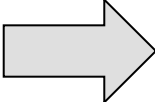
| Product | Date  | Price | Quantity |
|---------|-------|-------|----------|
| Bagel   | 10/21 | 1     | 20       |
| Bagel   | 10/25 | 1.50  | 20       |
| Banana  | 10/3  | 0.5   | 10       |
| Banana  | 10/10 | 1     | 10       |

## 2. Group by the attributes in the **GROUP BY**

```
SELECT product, SUM(price*quantity) AS TotalSales
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
```

| Product | Date  | Price | Quantity |
|---------|-------|-------|----------|
| Bagel   | 10/21 | 1     | 20       |
| Bagel   | 10/25 | 1.50  | 20       |
| Banana  | 10/3  | 0.5   | 10       |
| Banana  | 10/10 | 1     | 10       |

**GROUP BY**



| Product | Date  | Price | Quantity |
|---------|-------|-------|----------|
| Bagel   | 10/21 | 1     | 20       |
|         | 10/25 | 1.50  | 20       |
| Banana  | 10/3  | 0.5   | 10       |
|         | 10/10 | 1     | 10       |




### 3. Compute the **SELECT** clause: grouped attributes and aggregates

```
SELECT product, SUM(price*quantity) AS TotalSales  
FROM Purchase  
WHERE date > '10/1/2005'  
GROUP BY product
```

| Product | Date  | Price | Quantity |
|---------|-------|-------|----------|
| Bagel   | 10/21 | 1     | 20       |
|         | 10/25 | 1.50  | 20       |
| Banana  | 10/3  | 0.5   | 10       |
|         | 10/10 | 1     | 10       |

SELECT



| Product | TotalSales |
|---------|------------|
| Bagel   | 50         |
| Banana  | 15         |

# HAVING Clause

```
SELECT product, SUM(price*quantity)
FROM Purchase
WHERE date > '10/1/2005'
GROUP BY product
HAVING SUM(quantity) > 100
```

Same query as before, except that we consider only products that have more than 100 buyers

HAVING clauses contains conditions on **aggregates**

*Whereas WHERE clauses condition on **individual tuples**...*

# General form of Grouping and Aggregation

|          |                   |
|----------|-------------------|
| SELECT   | S                 |
| FROM     | $R_1, \dots, R_n$ |
| WHERE    | $C_1$             |
| GROUP BY | $a_1, \dots, a_k$ |
| HAVING   | $C_2$             |

*Why?*

- S = Can ONLY contain attributes  $a_1, \dots, a_k$  and/or aggregates over other attributes
- $C_1$  = is any condition on the attributes in  $R_1, \dots, R_n$
- $C_2$  = is any condition on the aggregate expressions

# General form of Grouping and Aggregation

|          |                   |
|----------|-------------------|
| SELECT   | S                 |
| FROM     | $R_1, \dots, R_n$ |
| WHERE    | $C_1$             |
| GROUP BY | $a_1, \dots, a_k$ |
| HAVING   | $C_2$             |

Evaluation steps:

1. Evaluate **FROM-WHERE**: apply condition  $C_1$  on the attributes in  $R_1, \dots, R_n$
2. **GROUP BY** the attributes  $a_1, \dots, a_k$
3. Apply condition  $C_2$  to each group (may need to compute aggregates)
4. Compute aggregates in S and return the result



# Example SQL


## Part I



# Example 1

## Sun Roof potential

from Satellite  
images

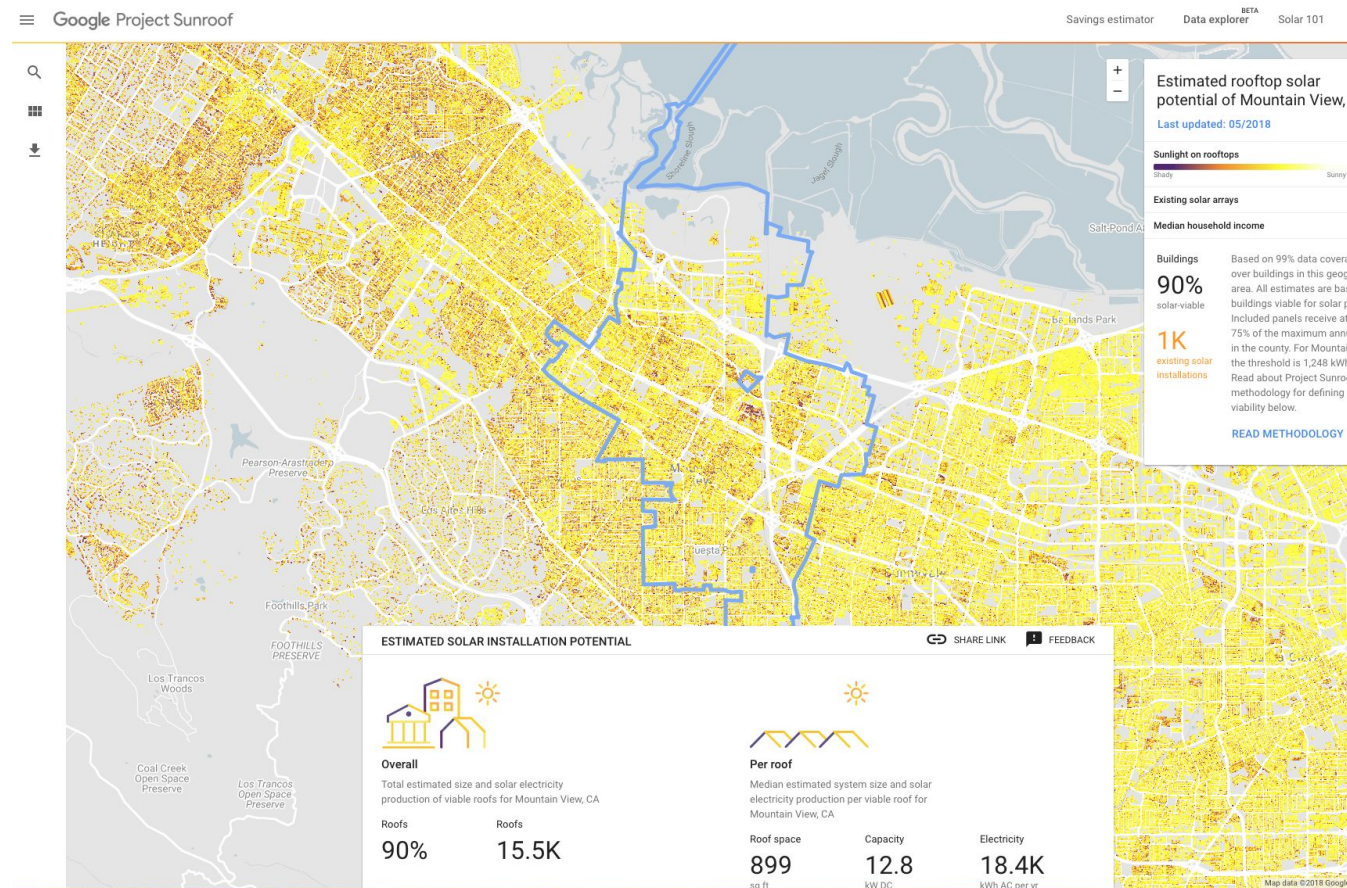




# Example 1

## SunRoof potential

### SunRoof explorer



## Public Dataset: Solar potential by postal code

### Example 1

| region_name | percent_covered   | kw_total | carbon_offset_metric_tons |
|-------------|-------------------|----------|---------------------------|
| 94043       | 97.79146031321109 | 215612.5 | 84929.00985071347         |
| 94041       | 99.05200433369447 | 56704.25 | 22189.34823862318         |

### Public dataset

## Public Dataset: USA.population by zip2010

| zipcode | population |
|---------|------------|
| 99776   | 124        |
| 38305   | 49808      |
| 37086   | 31513      |
| 41667   | 720        |
| 67001   | 1676       |



# Example 1

SunRoof

On BigQuery  
Public dataset

What is the solar potential of Mountain View, CA? [1]

🔗 Saved Query: MTV sunroof [edited] ?

```
1 #StandardSQL
2 SELECT
3     region_name,
4     percent_covered,
5     kw_total,
6     carbon_offset_metric_tons
7 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code`
8 WHERE
9     region_name = '94040'
10    OR region_name = '94041'
11    OR region_name = '94043'
```

Standard SQL Dialect ✕

Ctrl + Enter: run

RUN QUERY ▼

Save Query

Save View

Format Query

Schedule Query

Show Options

Query complete (1.6s elapsed, 346 KB processed)

Results

Details

Download as CSV

Download as JSON

Save

| Row | region_name | percent_covered   | kw_total | carbon_offset_metric_tons |  |
|-----|-------------|-------------------|----------|---------------------------|--|
| 1   | 94043       | 97.79146031321109 | 215612.5 | 84929.00985071347         |  |
| 2   | 94041       | 99.05200433369447 | 56704.25 | 22189.34823862318         |  |
| 3   | 94040       | 98.9440337909187  | 139745.5 | 55039.74974407879         |  |

## Example 2

SunRoof

Public dataset  
On BigQuery

How many metric tons of carbon would we offset, if building in communities with 100% coverage all had solar roofs?

🔗 Saved Query: CO2 offset in 100percent zips ?

```
1 #StandardSQL
2 SELECT
3     ROUND(SUM(s.carbon_offset_metric_tons),2) total_carbon_offset_possible_metric_tons
4 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code` s
5 JOIN `bigquery-public-data.census_bureau_usa.population_by_zip_2010` c
6 ON s.region_name = c.zipcode
7 WHERE
8     percent_covered = 100.0
9     AND c.population > 0
10
11
12
13
```

Standard SQL Dialect ✕

Ctrl + Enter: run query

RUN QUERY ▼

Save Query

Save View

Format Query

Schedule Query

Show Options

Query comments

Results

Details

Download as CSV

Download as JSON

Save as

| Row | total_carbon_offset_possible_metric_tons |
|-----|--|
| 1   | 3689508.33                               |

## Example 2

SunRoof

Public dataset  
On BigQuery

How many metric tons of carbon would we offset, per zipcode?

➞ Saved Query: CO2 offset in 100percent zips [edited] ?

```
1 #StandardSQL
2 SELECT
3   zipcode, ROUND(SUM(s.carbon_offset_metric_tons),2) total_carbon_offset_possible_metric_tons
4 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code` s
5 JOIN `bigquery-public-data.census_bureau_usa.population_by_zip_2010` c
6 ON s.region_name = c.zipcode
7 WHERE
8   percent_covered = 100.0
9   AND c.population > 0
10 GROUP BY c.zipcode
11
12
13
```

Standard SQL Dialect ✕

Ctrl + Enter: run c

RUN QUERY ▼

Save Query

Save View

Format Query

Schedule Query

Show Options

Query complete (2.5s elapsed, 23.5 MB processed)

Results

Details

Download as CSV

Download as JSON

Save a

| Row | zipcode | total_carbon_offset_possible_metric_tons |  |
|-----|---------|--|--|
| 1   | 35119   | 3417.26                                  |  |
| 2   | 10165   | 162.1                                    |  |
| 3   | 21810   | 6650.09                                  |  |
| 4   | 74078   | 61515.66                                 |  |
| 5   | 47876   | 5544.3                                   |  |
| 6   | 10170   | 831.22                                   |  |

# Example 2

SunRoof

How many metric tons of carbon would we offset, per zipcode

SQL Saved Query: CO2 offset in 100percent zips [edited] ?

Query Editor UDF

```
1 #StandardSQL
2 SELECT
3     zipcode, ROUND(SUM(s.carbon_offset_metric_tons),2) total_carbon_offset_possible_metric_tons
4 FROM `bigquery-public-data.sunroof_solar.solar_potential_by_postal_code` s
5 JOIN `bigquery-public-data.census_bureau_usa.population_by_zip_2010` c
6 ON s.region_name = c.zipcode
7 WHERE
8     percent_covered = 100.0
9     AND c.population > 0
10 GROUP BY c.zipcode
11 ORDER BY total_carbon_offset_possible_metric_tons
12 DESC
13
14
```

Standard SQL Dialect ✕

Ctrl + Enter: run query, Tab or Ctrl + Space

RUN QUERY

Save Query

Save View

Format Query

Schedule Query

Show Options

Query complete (2.8s elapsed, 23.5 MB processed)

Results

Details

Download as CSV

Download as JSON

Save as Table

Save to C

| Row | zipcode | total_carbon_offset_possible_metric_tons |
|-----|---------|--|
| 1   | 18503   | 715700.55                                |
| 2   | 44243   | 271861.55                                |
| 3   | 38677   | 266787.12                                |
| 4   | 96860   | 225850.35                                |
| 5   | 47809   | 141087.91                                |



# Query with SQL, universally over 'all' DBs

Reminder

Special  
Databases




SQL

```
= QUERY(T,  
  "SELECT c1, c2  
  FROM T  
  WHERE condition;")
```


```
SELECT c1, c2  
FROM T  
WHERE condition;
```

```
results =  
spark.SQL(  
  "SELECT c1, c2  
  FROM T  
  WHERE condition;")
```


DB



'Spreadsheets'



GCP BigQuery, AWS Redshift,  
MySQL, PostgreSQL, Oracle



Spark, Hadoop

Data

Data

100s of Scaling algorithms/systems? [Weeks 3..]

- Data layout? [Row vs columns...]
- Data structs? [Indexing...]